

# Relational Model and SQL

Concepts

Syntax

Basic Queries

Alvin Cheung  
Aditya Parameswaran  
Reading: R & G Chapter 5



# This Lecture

- The Relational Model
- SQL Basics



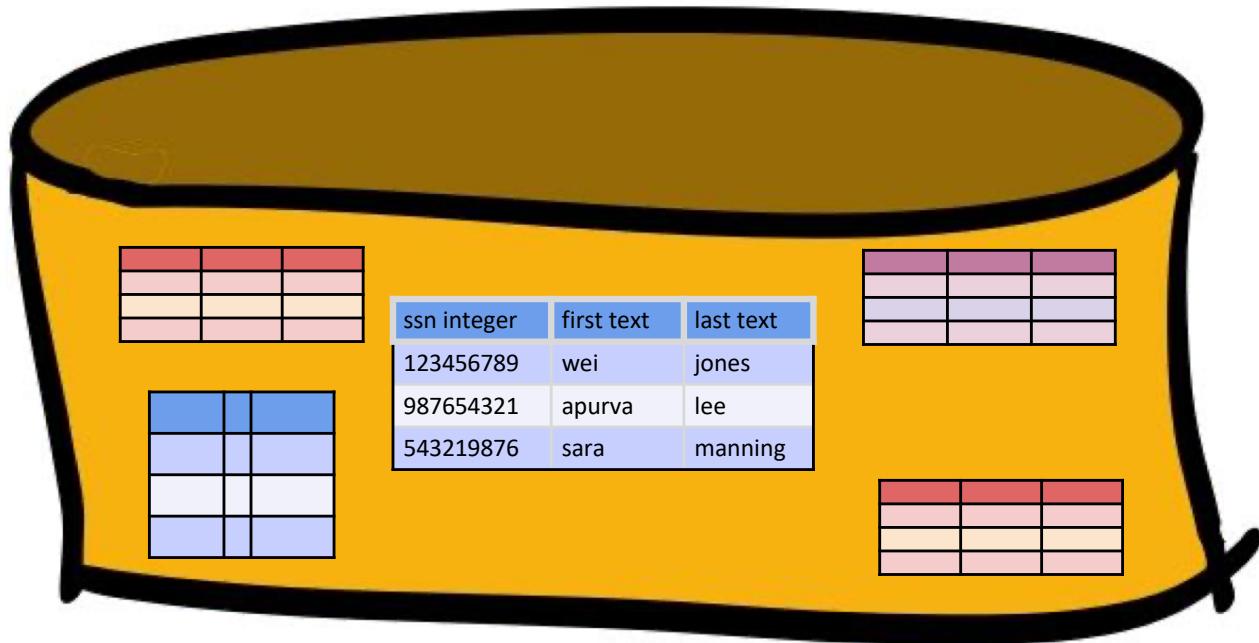
# ~~House~~ Zoom Rules



- Please turn on video if you feel comfortable
- Unmute for questions or comments
  - Raise hand / type in chat window works too

# Relational Terminology

- **Database:** Set of named Relations



# Relational Terminology, Pt 2.



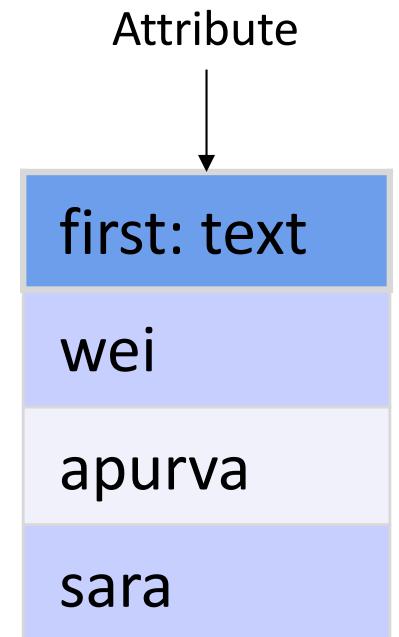
- **Database**: Set of named Relations
- **Relation** (aka *Table*):
  - **Schema**: description (“metadata”)
  - **Instance**: set of data satisfying the schema

ssn: integer	first: text	last: text
123456789	wei	jones
987654321	apurva	lee
543219876	sara	manning

# Relational Terminology, Pt. 3



- **Database:** Set of named Relations
- **Relation** (aka *Table*):
  - **Schema:** description (“metadata”)
  - **Instance:** set of data satisfying the schema
- **Attribute** (aka *Column, Field*)



# Relational Terminology, Pt. 4



- **Database**: Set of named Relations
- **Relation** (aka *Table*):
  - **Schema**: description (“metadata”)
  - **Instance**: set of data satisfying the schema
- **Attribute** (aka *Column, Field*)
- **Tuple** (aka *Record, Row*)

543219876	sara	manning
-----------	------	---------

A horizontal row of three cells, each containing a piece of data. A horizontal arrow points from the right side of the third cell towards the text "Tuple" located to its right.

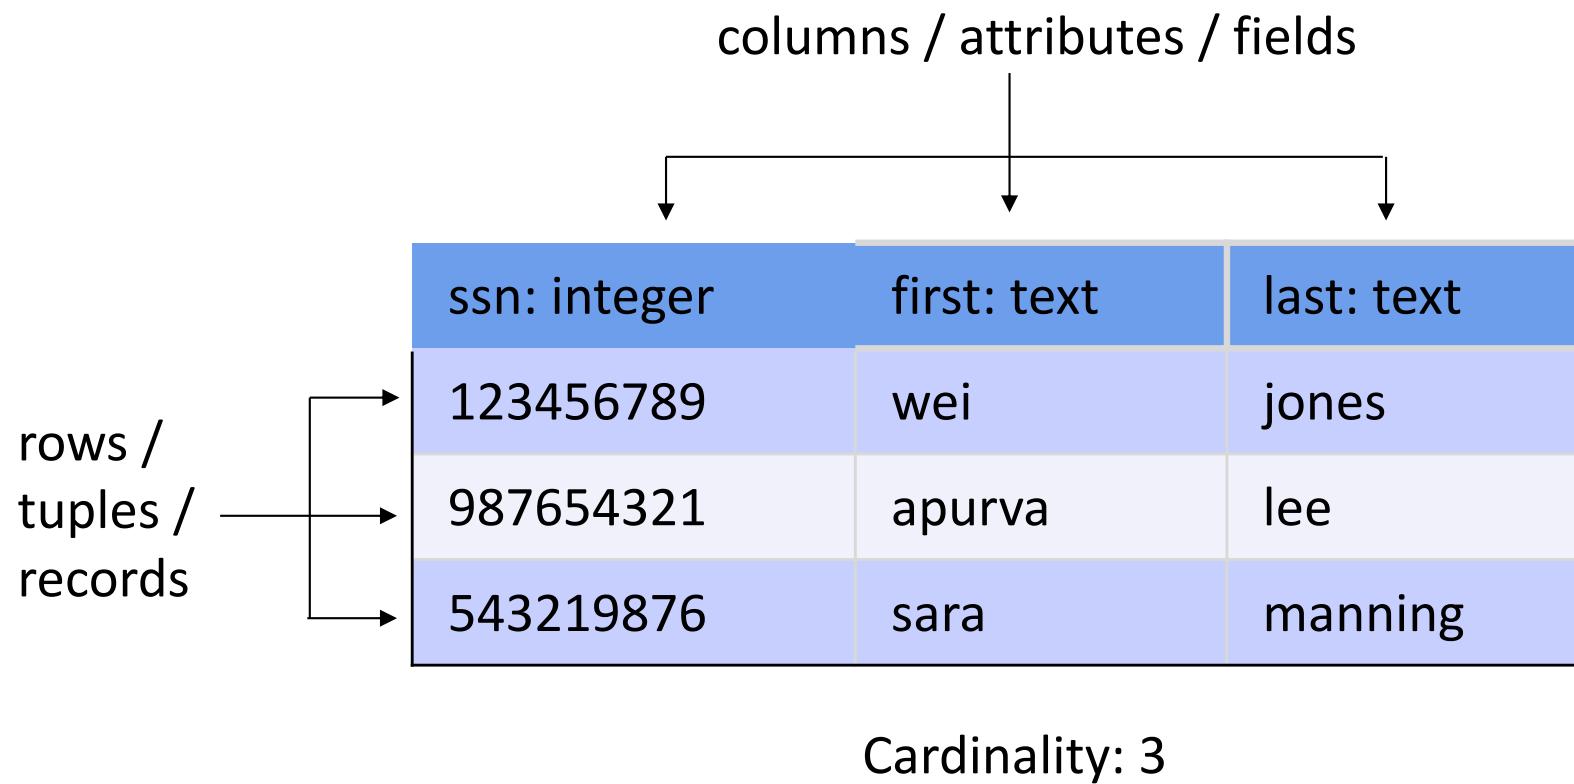
Tuple

# Relational Terminology, Pt. 5



- **Database**: Set of named Relations
- **Relation** (aka *Table*):
  - **Schema**: description (“metadata”)
  - **Instance**: set of data satisfying the schema
- **Attribute** (aka *Column, Field*)
- **Tuple** (aka *Record, Row*)
- **Cardinality**:
  - # of tuples in a relation

# To summarize



# Relational Tables



- *Schema* is fixed:
  - unique attribute names, *atomic* (aka *primitive*) types
- Tables are NOT ordered
  - they are sets or multisets (bags)
- Tables are FLAT
  - No nested attributes
- Tables DO NOT prescribe how they are implemented / stored on disk
  - This is called **physical data independence**

# Table Implementation



- How would you implement this?

cname	country	no_employees	for_profit
GizmoWorks	USA	20000	True
Canon	Japan	50000	True
Hitachi	Japan	30000	True
HappyCam	Canada	500	False

# Table Implementation



- How would you implement this?

cname	country	no_employees	for_profit
GizmoWorks	USA	20000	True
Canon	Japan	50000	True
Hitachi	Japan	30000	True
HappyCam	Canada	500	False

Row major: as an array of objects

GizmoWorks	Canon	Hitachi	HappyCam
USA	Japan	Japan	Canada
20000	50000	30000	500
True	True	True	False

# Table Implementation



- How would you implement this?

cname	country	no_employees	for_profit
GizmoWorks	USA	20000	True
Canon	Japan	50000	True
Hitachi	Japan	30000	True
HappyCam	Canada	500	False

Column major: as one array per attribute

GizmoWorks	Canon	Hitachi	HappyCam
USA	Japan	Japan	Canada
20000	50000	30000	500
True	True	True	False

# Table Implementation



- How would you implement this?

cname	country	no_employees	for_profit
GizmoWorks	USA	20000	True
Canon	Japan	50000	True
Hitachi	Japan	30000	True
HappyCam	Canada	500	False

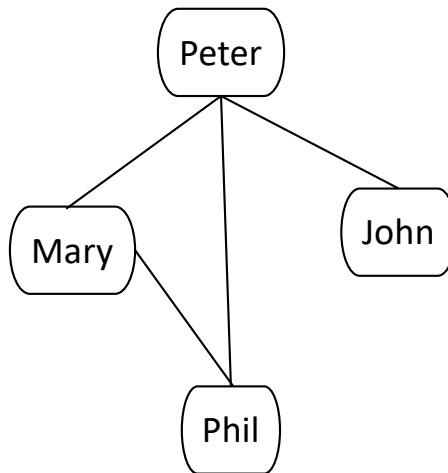
## Physical data independence

The logical definition of the data remains unchanged, even when we make changes to the actual implementation

# Relation is not the only data model



Example: storing FB friends



As a graph

OR

Person1: text	Person2: text	is_friend: int
Peter	John	1
John	Mary	0
Mary	Phil	1
Phil	Peter	1
...	...	...

As a relation

We will learn the tradeoffs of different  
data models in the semester

# Quick Check 1

- Why is this not a relation?

num: integer	street: text	zip: integer	
84	Maple Ave	54704	
22	High	Street	76425
75	Hearst Ave	94720	

## Quick Check 2

- Why is this not a relation?

num: integer	street: text	num: integer
84	Maple Ave	54704
22	High Street	76425
75	Hearst Ave	94720

# Quick Check 3

- Why is this not a relation?

first: text	last: text	addr: address
wei	jones	(84, 'Maple', 54704)
apurva	lee	(22, 'High', 76425)
sara	manning	(75, 'Hearst', 94720)

# First Normal Form



- All relations must be flat: we say that the relation is in *first normal form*

cname	country	no_employees	for_profit
Canon	Japan	50000	Y
Hitachi	Japan	30000	Y

# First Normal Form



- All relations must be flat: we say that the relation is in *first normal form*
- E.g., we want to add products manufactured by each company:

cname	country	no_employees	for_profit
Canon	Japan	50000	Y
Hitachi	Japan	30000	Y

# First Normal Form



- All relations must be flat: we say that the relation is in *first normal form*
- E.g., we want to add products manufactured by each company:

cname	country	no_employees	for_profit
Canon	Japan	50000	Y
Hitachi	Japan	30000	Y

cname	country	no_employees	for_profit	products									
Canon	Japan	50000	Y	<table border="1"><thead><tr><th>pname</th><th>price</th><th>category</th></tr></thead><tbody><tr><td>SingleTouch</td><td>149.99</td><td>Photography</td></tr><tr><td>Gadget</td><td>200</td><td>Toy</td></tr></tbody></table>	pname	price	category	SingleTouch	149.99	Photography	Gadget	200	Toy
pname	price	category											
SingleTouch	149.99	Photography											
Gadget	200	Toy											
Hitachi	Japan	30000	Y	<table border="1"><thead><tr><th>pname</th><th>price</th><th>category</th></tr></thead><tbody><tr><td>AC</td><td>300</td><td>Appliance</td></tr></tbody></table>	pname	price	category	AC	300	Appliance			
pname	price	category											
AC	300	Appliance											

# First Normal Form



- All relations must be flat: we say that the relation is in *first normal form*
- E.g., we want to add products manufactured by each company:

cname	country	no_employees	for_profit
Canon	Japan	50000	Y
Hitachi	Japan	30000	Y

Non-1NF!

cname	country	no_employees	for_profit	products									
Canon	Japan	50000	Y	<table border="1"><thead><tr><th>pname</th><th>price</th><th>category</th></tr></thead><tbody><tr><td>SingleTouch</td><td>149.99</td><td>Photography</td></tr><tr><td>Gadget</td><td>200</td><td>Toy</td></tr></tbody></table>	pname	price	category	SingleTouch	149.99	Photography	Gadget	200	Toy
pname	price	category											
SingleTouch	149.99	Photography											
Gadget	200	Toy											
<table border="1"><thead><tr><th>pname</th><th>price</th><th>category</th></tr></thead><tbody><tr><td>AC</td><td>300</td><td>Appliance</td></tr></tbody></table>	pname	price	category	AC	300	Appliance							
pname	price	category											
AC	300	Appliance											

# First Normal Form



Now it's in 1NF

Company

cname	country	no_employees	for_profit
Canon	Japan	50000	Y
Hitachi	Japan	30000	Y

Products

pname	price	category	manufacturer
SingleTouch	149.99	Photography	Canon
AC	300	Appliance	Hitachi
Gadget	200	Toy	Canon

# SQL Roots



- Developed @IBM Research in the 1970s
  - System R project
  - Vs. Berkeley's Quel language **INGRES**
- Commercialized/Popularized in the 1980s
  - IBM started the db2 product line
  - IBM beaten to market by a startup called Oracle

# SQL's Persistence



- Over 40 years old!
  - Not the only language for querying relations
- Questioned repeatedly
  - 90's: Object-Oriented DBMS (OQL, etc.)
  - 2000's: XML (Xquery, Xpath, XSLT)
  - 2010's: NoSQL & MapReduce
- SQL keeps re-emerging as the standard
  - Even Hadoop, Spark etc. mostly used via SQL
  - May not be perfect, but it is useful

# SQL Pros and Cons



- Declarative!
  - Say *what* you want, not *how* to get it
- Implemented widely
  - With varying levels of efficiency, completeness
- Constrained
  - Not targeted at Turing-complete tasks
- General-purpose for data computation and feature-rich
  - many years of added features
  - extensible: callouts to other languages, data sources

# SQL Language



- Two sublanguages:
  - DDL – Data Definition Language
    - Define and modify schema
  - DML – Data Manipulation Language
    - Queries can be written intuitively
- RDBMS responsible for efficient evaluation
  - Choose and run algorithms for declarative queries
    - Choice of algorithm must not affect query answer.

# Example Database



## Sailors

<u>sid</u>	sname	rating	age
1	Fred	7	22
2	Jim	2	39
3	Nancy	8	27

## Boats

<u>bid</u>	bname	color
101	Nina	red
102	Pinta	blue
103	Santa Maria	red

## Reserves

<u>sid</u>	bid	day
1	102	9/12/2015
2	102	9/13/2015

# The SQL DDL: Sailors



```
CREATE TABLE Sailors (
    sid INTEGER,
    sname CHAR(20),
    rating INTEGER,
    age FLOAT)
```

<b><u>sid</u></b>	<b>sname</b>	<b>rating</b>	<b>age</b>
1	Fred	7	22
2	Jim	2	39
3	Nancy	8	27

# The SQL DDL: Sailors, Pt. 2



```
CREATE TABLE Sailors (
    sid INTEGER,
    sname CHAR(20),
    rating INTEGER,
    age FLOAT
    PRIMARY KEY (sid));
```

<b><u>sid</u></b>	<b>sname</b>	<b>rating</b>	<b>age</b>
1	Fred	7	22
2	Jim	2	39
3	Nancy	8	27

# The SQL DDL: Primary Keys



```
CREATE TABLE Sailors (
    sid INTEGER,
    sname CHAR(20),
    rating INTEGER,
    age FLOAT,
    PRIMARY KEY (sid))
```

<b><u>sid</u></b>	<b>sname</b>	<b>rating</b>	<b>age</b>
1	Fred	7	22
2	Jim	2	39
3	Nancy	8	27

- Primary Key column(s)
  - Provides a unique “lookup key” for the relation
  - Cannot have any duplicate values
  - Can be made up of >1 column
    - E.g. (firstname, lastname)

# The SQL DDL: Boats



```
CREATE TABLE Sailors (
    sid INTEGER,
    sname CHAR(20),
    rating INTEGER,
    age FLOAT,
    PRIMARY KEY (sid));
```

<b><u>sid</u></b>	<b>sname</b>	<b>rating</b>	<b>age</b>
1	Fred	7	22
2	Jim	2	39
3	Nancy	8	27

```
CREATE TABLE Boats (
    bid INTEGER,
    bname CHAR (20),
    color CHAR(10),
    PRIMARY KEY (bid));
```

<b><u>bid</u></b>	<b>bname</b>	<b>color</b>
101	Nina	red
102	Pinta	blue
103	Santa Maria	red

# The SQL DDL: Reserves



```
CREATE TABLE Sailors (
    sid INTEGER,
    sname CHAR(20),
    rating INTEGER,
    age FLOAT,
    PRIMARY KEY (sid));
```

```
CREATE TABLE Boats (
    bid INTEGER,
    bname CHAR(20),
    color CHAR(10),
    PRIMARY KEY (bid));
```

```
CREATE TABLE Reserves (
    sid INTEGER,
    bid INTEGER,
    day DATE,
    PRIMARY KEY (sid, bid, day));
```

<u>sid</u>	<u>sname</u>	<u>rating</u>	<u>age</u>
1	Fred	7	22
2	Jim	2	39
3	Nancy	8	27

<u>bid</u>	<u>bname</u>	<u>color</u>
101	Nina	red
102	Pinta	blue
103	Santa Maria	red

<u>sid</u>	<u>bid</u>	<u>day</u>
1	102	9/12
2	102	9/13

# The SQL DDL: Reserves Pt. 2



```
CREATE TABLE Sailors (
    sid INTEGER,
    sname CHAR(20),
    rating INTEGER,
    age FLOAT,
    PRIMARY KEY (sid));
```

```
CREATE TABLE Boats (
    bid INTEGER,
    bname CHAR(20),
    color CHAR(10),
    PRIMARY KEY (bid));
```

```
CREATE TABLE Reserves (
    sid INTEGER,
    bid INTEGER,
    day DATE,
    PRIMARY KEY (sid, bid, day),
    FOREIGN KEY (sid) REFERENCES Sailors,
```

<u>sid</u>	<u>sname</u>	<u>rating</u>	<u>age</u>
1	Fred	7	22
2	Jim	2	39
3	Nancy	8	27

<u>bid</u>	<u>bname</u>	<u>color</u>
101	Nina	red
102	Pinta	blue
103	Santa Maria	red

<u>sid</u>	<u>bid</u>	<u>day</u>
1	102	9/12
2	102	9/13

# The SQL DDL: Foreign Keys



```
CREATE TABLE Sailors (
    sid INTEGER,
    sname CHAR(20),
    rating INTEGER,
    age FLOAT,
    PRIMARY KEY (sid));
```

```
CREATE TABLE Boats (
    bid INTEGER,
    bname CHAR(20),
    color CHAR(10),
    PRIMARY KEY (bid));
```

```
CREATE TABLE Reserves (
    sid INTEGER,
    bid INTEGER,
    day DATE,
    PRIMARY KEY (sid, bid, day),
    FOREIGN KEY (sid) REFERENCES Sailors,
    FOREIGN KEY (bid) REFERENCES Boats);
```

sid	sname	rating	age
1	Fred	7	22
2	Jim	2	39
3	Nancy	8	27

bid	bname	color
101	Nina	red
102	Pinta	blue
103	Santa Maria	red

sid	bid	day
1	102	9/12
2	102	9/13

# The SQL DDL: Foreign Keys Pt. 2



- Foreign key references a table
  - Via the primary key of that table
- Doesn't need to have the same name as the referenced primary key

```
CREATE TABLE Reserves (
    sid INTEGER,
    bid INTEGER,
    day DATE,
    PRIMARY KEY (sid, bid, day),
    FOREIGN KEY (sid) REFERENCES Sailors,
    FOREIGN KEY (bid) REFERENCES Boats);
```

sid	sname	rating	age
1	Fred	7	22
2	Jim	2	39
3	Nancy	8	27

bid	bname	color
101	Nina	red
102	Pinta	blue
103	Santa Maria	red

sid	bid	day
1	102	9/12
2	102	9/13

# The SQL DML



- Find all 27-year-old sailors:

```
SELECT *
FROM Sailors AS S
WHERE S.age=27;
```

- To find just names and rating, replace the first line to:

```
SELECT S.sname,
S.rating
```

**Sailors**

<b><u>sid</u></b>	<b>sname</b>	<b>rating</b>	<b>age</b>
1	Fred	7	22
2	Jim	2	39
3	Nancy	8	27

# Basic Single-Table Queries



- **SELECT** [*DISTINCT*] <column expression list>  
**FROM** <single table>  
[**WHERE** <predicate>]
- In this simple version:
  - Produce all tuples in the table that satisfy the predicate
  - Output the expressions in the SELECT list
  - Expression can be a column reference, or an arithmetic expression over column refs

# Distinct and Alias



```
SELECT DISTINCT S.name, S.gpa  
FROM students [AS] S  
WHERE S.dept = 'CS'
```

- Return all unique (name, GPA) pairs from students
- DISTINCT specifies removal of duplicate rows before output
- Can refer to the students table as S, this is called an *alias*

# Ordering



- ```
SELECT S.name, S.gpa, S.age*2 AS a2
  FROM Students S
 WHERE S.dept = 'CS'
 ORDER BY S.gpa, S.name, a2;
```
- ORDER BY clause specifies output to be sorted
  - Numeric ordering for “number-like” attributes (int, real, etc)
  - **Lexicographic** ordering otherwise (!!)(varchar, blob, etc)
- Obviously must refer to columns in the output
  - Note the AS clause for naming output columns!

# Ordering



- ```
SELECT S.name, S.gpa, S.age*2 AS a2
  FROM Students S
 WHERE S.dept = 'CS'
 ORDER BY S.gpa DESC, S.name ASC, a2;
```
- Ascending order by default, but can be overridden
  - DESC flag for descending, ASC for ascending
  - Can mix and match, lexicographically

# Setting limits



- ```
SELECT S.name, S.gpa, S.age*2 AS a2
  FROM Students S
 WHERE S.dept = 'CS'
 ORDER BY S.gpa DESC, S.name ASC, a2;
 LIMIT 3 ;
```
- Only produces the first <integer> output rows
- Typically used with ORDER BY
  - Otherwise the output is *non-deterministic*
  - Not a “pure” declarative construct in that case – output set depends on algorithm for query processing

# Aggregates



- ```
SELECT [DISTINCT] AVG(S.gpa)
  FROM Students S
 WHERE S.dept = 'CS'
```
- Before producing output, compute a summary (aka an *aggregate*) of some arithmetic expression
- Produces 1 row of output
  - with one column in this case
- Other aggregates: SUM, COUNT, MAX, MIN (and others)

# DISTINCT Aggregates

Are these the same or different?

```
SELECT COUNT(DISTINCT S.name)
FROM Students S
WHERE S.dept = 'CS';
```

```
SELECT DISTINCT COUNT(S.name)
FROM Students S
WHERE S.dept = 'CS';
```



# GROUP BY



```
SELECT [DISTINCT] AVG(S.gpa), S.dept  
FROM Students S  
GROUP BY S.dept
```

- Partition table into groups with same GROUP BY column values
  - Can group by a list of columns
- Produce an aggregate result per group
  - Cardinality of output = # of distinct group values
- Note: only grouping columns or aggregated values can appear in the SELECT list

# Need to be Careful with GROUP BY...

```
SELECT product,  
       max(quantity)  
FROM   Purchase  
GROUP  BY product
```

Purchase

Product	Price	Quantity
Bagel	3	20
Bagel	1.50	20
Banana	0.5	50
Banana	2	10
Banana	4	10



# Need to be Careful with GROUP BY...



```
SELECT product,  
       max(quantity)  
FROM   Purchase  
GROUP  BY product
```

Purchase

Product	Price	Quantity
Bagel	3	20
Bagel	1.50	20
Banana	0.5	50
Banana	2	10
Banana	4	10

Product	Max(quantity)
Bagel	20
Banana	50

# Need to be Careful with GROUP BY...



```
SELECT product,  
       max(quantity)  
FROM   Purchase  
GROUP  BY product
```

```
SELECT product, quantity  
FROM   Purchase  
GROUP  BY product  
-- what does this mean?
```

Purchase

Product	Price	Quantity
Bagel	3	20
Bagel	1.50	20
Banana	0.5	50
Banana	2	10
Banana	4	10

Product	Max(quantity)
Bagel	20
Banana	50

# Need to be Careful with GROUP BY...



```
SELECT product,  
       max(quantity)  
FROM   Purchase  
GROUP  BY product
```

```
SELECT product, quantity  
FROM   Purchase  
GROUP  BY product  
-- what does this mean?
```

Purchase

Product	Price	Quantity
Bagel	3	20
Bagel	1.50	20
Banana	0.5	50
Banana	2	10
Banana	4	10

Product	Max(quantity)
Bagel	20
Banana	50

Product	Quantity
Bagel	20
Banana	??

# Need to be Careful with GROUP BY...



```
SELECT product,  
       max(quantity)  
FROM   Purchase  
GROUP  BY product
```

```
SELECT product, quantity  
FROM   Purchase  
GROUP  BY product  
-- what does this mean?
```

Purchase

Product	Price	Quantity
Bagel	3	20
Bagel	1.50	20
Banana	0.5	50
Banana	2	10
Banana	4	10

Product	Max(quantity)
Bagel	20
Banana	50

Product	Quantity
Bagel	20
Banana	??



# Need to be careful with GROUP BY...

Everything in SELECT must be either a GROUP-BY attribute, or an aggregate

```
SELECT product,  
       max(quantity)  
FROM Purchase  
GROUP BY product
```

```
SELECT product, quantity  
FROM Purchase  
GROUP BY product  
-- what does this mean?
```

Purchase

Product	Price	Quantity
Bagel	3	20
Bagel	1.50	20
Banana	0.5	50
Banana	2	10
Banana	4	10

Product	Max(quantity)
Bagel	20
Banana	50

Product	Quantity
Bagel	20
Banana	??



# HAVING

```
SELECT [DISTINCT] AVG(S.gpa), S.dept  
FROM Students S  
GROUP BY S.dept  
HAVING COUNT(*) > 2
```



- The HAVING predicate filters groups
- HAVING is applied *after* grouping and aggregation
  - Hence can contain anything that could go in the SELECT list
  - i.e., aggs or GROUP BY columns
- HAVING can only be used in aggregate queries
- It's an optional clause

# SQL DML: General Single-Table Queries



- **SELECT [DISTINCT] <column expression list>  
FROM <single table>  
[WHERE <predicate>]  
[GROUP BY <column list>  
[HAVING <predicate> ]  
[ORDER BY <column list>]  
[LIMIT <integer>];**

# Summary



- Many query languages available for the relational data model
  - SQL is one of them that we will focus in this class
- Modern SQL extends set-based relational model
  - some extra goodies for duplicate row (bags), non-atomic types...
- Typically, many ways to write a query
  - DBMS figures out a fast way to execute a query, regardless of how it is written